

# DATA SCIENCE II: Machine Learning MTH 9899 Baruch College

## Lecture 2: Machine Learning

Adrian Sisser

April 5, 2017

# Outline

- 1 Linear Models
  - Subset Selection
- 2 K-Means, Clustering, and EM
  - The K-Means Algorithm
  - The EM Algorithm
- 3 Training Neural Networks

# Outline

- 1 Linear Models
  - Subset Selection
- 2 K-Means, Clustering, and EM
  - The K-Means Algorithm
  - The EM Algorithm
- 3 Training Neural Networks

# Linear Methods

- One hard problem with Regression, is to choose which variables are relevant to your model.
- You might have dozens of predictors, most of which might be irrelevant.
- This problem is much harder in the context of low  $R^2$  and correlated predictors.

One technique is “Best Subset” regression, which finds the optimal subset of size  $k$  regressors for each value  $k$ . Computationally, this is very expensive.

# Forward Stepwise Regression

- We look for the best variable from the remaining ones at each stage and add it into the regression,
- This again gives us  $k$  models
- This is a greedy algorithm that might not be best.
- It will ALWAYS be worse than the corresponding best subset, but has lower variance.
- Caveat: Imagine you have 2 highly correlated  $x$  variables that predict  $y$ , but both are measured with significant noise. What will Forward Stepwise do? What would you prefer?

# Backward Stepwise Regression

- We start off with all variables and remove the 'worst' one at each stage.
- Same Caveat as Forward Stepwise Regression

# Model Selection

- AIC - Akaike Information Criterion. This can be shown to be asymptotically equivalent to  $N$ -Fold Cross-Validation [Stone, 1977].

$$\text{AIC} = 2k - 2 \ln(\mathcal{L})$$

- BIC - Bayesian Information Criterion

$$\text{BIC} = \ln(n)k - 2 \ln(\mathcal{L})$$

# Lasso Regression

- Lasso is similar to Ridge Regression, except we use the L1 penalty:

$$\min_{\beta^L} \|Y - X\hat{\beta}^L\| + \lambda\|\beta^L\|_1$$

- Can we compute an analytical solution to this?




# LARS

LARS - Least Angle Regression, is a technique with an intuitive geometric explanation that leads to a very efficient implementation of LASSO. For more details, the class text, Elements of Statistical Learning provides the best intuition of how it works.

# Lasso vs Ridge Regression

- Lasso works well for Feature Selection
- Ridge works well for Correlation features

Elements of Statistical Learning (2nd Ed.) ©Hastie, Tibshirani & Friedman 2009 Chap 3



# Elastic Net Regression

- Let's get the best of both worlds - we can use an L1 and L2 Penalty:

$$\min_{\beta^{EN}} \|Y - X\hat{\beta}^{EN}\| + \lambda_1 \|\beta^{EN}\|_1 + \lambda_2 \|\beta^{EN}\|_2$$

# Outline

- 1 Linear Models
  - Subset Selection
- 2 K-Means, Clustering, and EM
  - The K-Means Algorithm
  - The EM Algorithm
- 3 Training Neural Networks

# K-Means Introduction


K-Means is a method of classifying points into 'groups' or 'clusters', based on their 'proximity'. For traditional k-means, the proximity measure is Euclidean distance, ie  $\| \cdot \|_2$ . If we want to form  $K$  clusters, we minimize as follows:

$$\arg \min_S \sum_{i=0}^K \sum_{x \in S_i} \|x - C_i\|$$

where  $C_i$  is the geometric center of all of the points belonging to  $S_i$ , the set of all points in cluster  $i$ . This is equivalent to assuming the points are normally distributed around each center.


# K-Means Introduction

Here is sample of clusters in 2 dimensions.




# K-Means Introduction

Here is sample of clusters in 2 dimensions.



# K-Means Introduction

Here is sample of clusters in 2 dimensions.





# K-Means Algorithm

So how do we do this? We need an algorithm to solve the minimization problem from earlier:

$$\arg \min_S \sum_{i=0}^K \sum_{x \in S_i} \|x - C_i\|$$

Solving this problem directly isn't tractable - in fact, it's *NP*-hard for almost all cases.

# KMeans Algorithm

---

## Algorithm 1 KMeans Algorithm (Lloyd's Algorithm)

---

**Require:**  $N > K > 1$

$centers \leftarrow$  select  $K$  random entries from  $points$

**repeat**

**for**  $i < N$  **do**

$assigned\_centers[i] \leftarrow \text{find\_nearest\_center}(points[i])$

**end for**

**for**  $i < K$  **do**

$centroids[i] \leftarrow \text{find\_centroid}(i)$

**end for**

**until**  $assigned\_centers$  does not change


---

# Initialization

So how do we *randomly* assign the initial clusters? There are a few popular choices:


- Choose  $K$  random points from the initial list (Forgy Method).
- The Random partition method assigns each point a cluster at random, then calculates the centroids.

# An example




0 iterations

# An example




1 iterations

# An example




2 iterations

# An example



5 iterations


# An example



10 iterations




# An example




20 iterations

# An example





50 iterations

# An example



Our original data

# An example



Actual Centers

( 2.0, 5.0)

→

( 2.0, 0.0)

→

(-1.0, -1.0)

→

Calculated Centers

( 1.9, 4.7)

( 1.9, 0.2)

(-0.8, -1.1)

# Notes

- The algorithms discussed will only find a **LOCAL** minimum. To be sure we're getting a near-optimal solution, we should repeat this with different starting centroids.
- How do we know how many clusters,  $K$ , to look for? Adding more clusters will always improve the metrics.

# GMeans

G-Means offers a way for us to intuit  $K$ :

---

## Algorithm 2 GMeans Algorithm

---

```

 $K \leftarrow 0$ 
repeat
     $K \leftarrow K + 1$ 
     $centers \leftarrow \text{KMeans}(points, K)$ 
until  $(points - centers) \sim \mathcal{N}$ 
    
```

---

# Outline

- 1 Linear Models
  - Subset Selection
- 2 K-Means, Clustering, and EM
  - The K-Means Algorithm
  - The EM Algorithm
- 3 Training Neural Networks

# The EM Algorithm

“Expectation-Maximization” is a generic algorithm for estimating MLE parameters. The derivation is complex, and we will go through it quickly here. An excellent reference is Andrew Ng’s ML Notes.

$$X = \{x_0, x_1, \dots, x_{n-2}, x_{n-1}\}$$

$$Z = \{z_0, z_1, \dots, z_{n-2}, z_{n-1}\} \text{ \# These are our latent variables}$$

$$\mathcal{L}(\theta|X, Z) = \prod_{i=0}^N P(x_i; \theta)$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log P(x_i; \theta)$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \sum_{j=0}^K P(x_i, z_j; \theta)$$



# The EM Algorithm

Let's define  $Q_i$  as a probability distribution of  $z_i$ . Now we can say:

$$\begin{aligned}\ell(\theta|X, Z) &= \sum_{i=0}^N \log \sum_{j=0}^K P(x_i, z_i; \theta) \\ \ell(\theta|X, Z) &= \sum_{i=0}^N \log \sum_{j=0}^K Q_i(z_j) \frac{P(x_i, z_i; \theta)}{Q_i(z_j)}\end{aligned}$$

# The EM Algorithm

Let's define  $Q_i$  as a probability distribution of  $z_i$ . Now we can say:

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \sum_{j=0}^K P(x_i, z_i; \theta)$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \sum_{j=0}^K Q_i(z_j) \frac{P(x_i, z_i; \theta)}{Q_i(z_j)}$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \mathbb{E}_{z_j \sim Q_i(z_j)} \frac{P(x_i, z_i; \theta)}{Q_i(z_j)}$$

# The EM Algorithm

Let's define  $Q_i$  as a probability distribution of  $z_i$ . Now we can say:

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \sum_{j=0}^K P(x_i, z_j; \theta)$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \sum_{j=0}^K Q_i(z_j) \frac{P(x_i, z_j; \theta)}{Q_i(z_j)}$$

$$\ell(\theta|X, Z) = \sum_{i=0}^N \log \mathbb{E}_{z_j \sim Q_i(z_j)} \frac{P(x_i, z_j; \theta)}{Q_i(z_j)}$$

$$\ell(\theta|X, Z) \geq \sum_{i=0}^N \mathbb{E}_{z_j \sim Q_i(z_j)} \log \frac{P(x_i, z_j; \theta)}{Q_i(z_j)}$$

By Jensen's  
Inequality

$$\ell(\theta|X, Z) \geq \sum_{i=0}^N \sum_{j=0}^K Q_i(z_j) \log \frac{P(x_i, z_j; \theta)}{Q_i(z_j)}$$

# The EM Algorithm

The next steps are tricky (again, refer to Andrew Ng's ML Notes for more details). We said that  $Q_i$  was a PDF for  $z_i$ , so let's choose a good one:

$$\begin{aligned} Q_i(z_i) &= \frac{P(x_i, z_i; \theta)}{\sum_j P(x_i, z_j; \theta)} \\ &= P(z_i | x_i; \theta) \end{aligned}$$

Now, we're ready to look at the algorithm itself.

# The EM Algorithm

---

## Algorithm 3 EM Algorithm

---

$\theta^0 =$  initial guess

$m \leftarrow 1$

**repeat**

$$Q_i^m = P(z_i | x_i; \theta^m)$$

$$\theta^{m+1} = \arg \max_{\theta} \sum_{i=0}^N \sum_{j=0}^K Q_i^m(z_j) \log \frac{P(x_i, z_i; \theta^{m+1})}{Q_i^m(z_j)}$$

$m \leftarrow m + 1$

**until** convergence of  $\ell$

---

**Take careful note of  $\theta$  in the MLE step.** Proof of convergence can be found in the Ng reference mentioned above.

# An EM Application: Soft KMeans

Let's look at this in the context of a 'soft' KMeans Algorithm with 2 clusters. This means that instead of assuming each point is in a given cluster,  $C$ , we'll assign a probability that it's in each cluster. Here's our setup:

$$\begin{aligned}X &= \{x_0, x_1, \dots, x_n\} \\Z &= \{z_0, z_1, \dots, z_n\} \\ \theta &= \{\mu_0, \sigma_0^2, \mu_1, \sigma_1^2, \pi_0, \pi_1\}\end{aligned}$$

# Soft KMeans: The “E” Step

$$\begin{aligned} Q_i(z_j) &= \frac{P(x_i, z_j; \theta)}{\sum_m P(x_i, z_m; \theta)} \\ &= \frac{\phi_j(x_i; \theta)}{\sum_m \phi_m(x_i; \theta)} \end{aligned}$$

$Q_i(z_j)$  is the probability that point  $i$  belongs to  $C_j$ . Since we don't make a hard assignment to any cluster, this is why we call this a 'Soft K-Means' algorithm.

# Soft KMeans: The “M” Step

To make notation simpler, now that we’ve done the “E” step, we’ll say  $w_{i,j}$  is the probability that point  $i$  is in  $C_j$ . The “M” step is:

$$\arg \max_{\theta} \sum_{i=0}^N \sum_{j=0}^K w_{i,j} \log \frac{P(x_i, z_j; \theta)}{w_{i,j}}$$

$$\arg \max_{\theta} \sum_{i=0}^N \sum_{j=0}^K w_{i,j} \log \frac{P(x_i|z_j; \theta)P(z_j)}{w_{i,j}}$$

$$\arg \max_{\theta} \sum_{i=0}^N \sum_{j=0}^K w_{i,j} \log \frac{\phi_{j;\theta}(x_i)\pi_j}{w_{i,j}}$$




# Soft KMeans: The “M” Step

$$\arg \max_{\theta} \sum_{i=0}^N \sum_{j=0}^K w_{i,j} \log \frac{\phi_{j;\theta}(x_i) \pi_j}{w_{i,j}}$$

If we take our function from before, and take some derivatives, we get very simple update rules:


$$\begin{aligned} \mu_j &= \frac{\sum_{i=0}^N w_{ij} x_i}{\sum_{i=0}^N w_{ij}} \\ \pi_j &= \frac{1}{N} \sum_i w_{ij} \\ \sigma_j^2 &= \frac{\sum_{i=0}^N w_{ij} \|x_i - \mu_j\|_2^2}{\sum_{i=0}^N w_{ij}} \end{aligned}$$

# An example




0 iterations

# An example




1 iterations

# An example




2 iterations

# An example




5 iterations

# An example




10 iterations

# An example



20 iterations


# An example



50 iterations





# An example



Our original data

# An example



Actual Centers

( 2.0, 5.0)

( 2.0, 0.0)

(-1.0, -1.0)

→

Calculated Centers

( 1.9, 4.7)

→

( 1.8, 0.2)

→

(-0.7, -1.0)

# Why Training is Important

It can be show that a sufficiently large NN can learn any function. The hardest part is training the weights in the network. Why is it hard?

- The number of weights in a network is huge. Connections between an  $N$  and  $M$  neuron layer create a total of  $\mathcal{O}(NM)$  connections.
- A training algorithm on a deep network based on gradients suffers from the *Vanishing/Exploding Gradient Problem*.

*In summary, if we design a network that can learn anything, it's REALLY hard to make it learn what we want.*

# Basic Training


- When we talk about training a NN, we're talking about finding weights and biases for the neurons.
- We want to adjust the weights, such that we reduce the error.
- To do this, we'll follow the gradient of the error with respect to the weights.

# Activation Functions

Before we can talk about training, we need to talk about Activation Functions. For now, let's talk about sigmoid - a very common and simple activation function. Later, we'll see why we need this.

$$\varphi(z) = \frac{1}{1 + e^{-z}}$$

$$\frac{d\varphi}{dz} = \varphi(z)(1 - \varphi(z))$$




# Backpropagation

Based on the last lecture, let's assume that a neuron can be modeled as:

$$n_j = A\left(\sum_i w_{ij}n_i + b_j\right)$$

And a simple network topology:



# Backpropagation

We're going to come up with a learning rule. The goal is to reduce the total error. For a simple regression problem:

$$E = \sum_i \frac{1}{2} (\hat{y}_i - y_i)^2$$

Now that we know the error, let's calculate the derivative with respect to  $\hat{y}$ :

$$\frac{\partial E}{\partial \hat{y}_i} = (\hat{y}_i - y_i)$$

# Backpropagation

Now, we need to propagate the errors backwards, through the NN. First, let's setup some notations:

- $x_i$  Row  $i$  of the feature matrix,  $X$ .
- $n_{l,i}$  is the value of the output of neuron  $i$  in layer  $l$
- $\hat{y}_i$  The prediction for sample  $i$ , ie the output of the neuron
- $w_{l,ij}$  The weights from the ( $l$ ) layer, neuron  $i$ , to the  $l + 1$  layer, neuron  $j$ .
- $b_{l,j}$  The bias of of neuron  $j$  in layer  $l$



# Backpropagation

Let's look at the error term for the hidden layer in the topology we talked about. We'll say  $z = \sum_i w_{h,ij} n_{h,i} + b_j$  for notation.

$$\begin{aligned}\frac{\partial E}{\partial w_{h,i0}} &= \frac{\partial E}{\partial n_{o,0}} \frac{\partial n_{o,0}}{\partial w_{h,i0}} \\ \frac{\partial E}{\partial w_{h,i0}} &= \frac{\partial E}{\partial n_{o,0}} \frac{\partial A(z)}{\partial w_{h,i0}} \\ \frac{\partial E}{\partial w_{h,i0}} &= \frac{\partial E}{\partial n_{o,0}} \frac{\partial \varphi(z)}{\partial z} \frac{\partial z}{\partial w_{h,i0}} \\ \frac{\partial E}{\partial w_{h,i0}} &= (\hat{y}_i - y_i) \varphi(z) (1 - \varphi(z)) n_{h,i}\end{aligned}$$

# Backpropagation

Now we can finally turn this into a learning rule.

$$w'_{l,ij} = w_{l,ij} - \eta \frac{\partial E}{\partial w_{l,ij}}$$